

Genestrip: Space-efficient k -mer counting, read classification for selected groups of organisms

Daniel Pfeifer^{1*}, Markus Graf¹ and Clas Rurik²

^{1*}IT Faculty, Heilbronn University, Max-Planck-Str. 39, Heilbronn, 74081, Baden-Württemberg, Germany.

²codeclap, Kokhaven 20, Røddekro, 6230 Denmark.

*Corresponding author(s). E-mail(s): daniel.pfeifer@hs-heilbronn.de;
Contributing authors: markus.graf@hs-heilbronn.de; cru@codeclap.cc;

Abstract

Background: The consumption of main memory resources is a significant burden in k -mer-based metagenomic analysis when creating related databases but also when performing (unique) k -mer-counting and read classification. Genestrip addresses this issue by focusing on small but freely configurable groups of organisms. Regarding the selected organisms, Genestrip produces k -mer databases and results comparable to those of KrakenUniq but at a fraction of its required memory resources. Our tool ensures that during database generation, the most suitable lowest common ancestor taxon is assigned for each stored k -mer by *also considering genomes of organisms whose k -mers are not included in the database*. This enables read analysis with high precision and recall for the organisms of interest.

Results: We assessed the correctness, usefulness and performance of Genestrip in different contexts and showed that it indeed ascertains high quality read classifications for organisms whose genomes were included in a corresponding database. Our example databases comprise millions to a few billions of k -mers covering a dozen to a few thousands of species and lend themselves to usage in tick surveillance, medical diagnostics or agriculture. All databases were generated on a regular PC within hours, and related analysis performance was competitive to highly favorable. The deliberate focus on a particular set of genera or species allowed for more genomes to be included from related organisms while the resulting databases remained small. We exemplify, that such small but deep databases tend to improve recall during read classification while sustaining high precision.

Conclusions: Due to Genestrip's particular way of updating the k -mers' lowest common ancestor taxa, both database creation and fastq file analysis can be

Rework
to
address
false
positives
and
 k -
mer
compression,
less
repeat
in the
results
section...

realized with little memory and with favorable runtimes as well as high classification quality. So both, database creation *and* read classification may be performed even on regular PCs. Genestrip’s qualities empower users to flexibly design, build and use small k -mer databases for their own needs with potentially deep genomic coverage.

Keywords: Metagenomics, k -mer counting, read classification, small k -mer database, fastq file filtering, high accuracy, high performance

1 Background

1.1 Introduction

With the availability of low-cost and small-scale sequencing devices, metagenomic analysis will likely replace many established technologies for the detection and surveillance of specific groups of microorganisms. Related applications may affect various areas such as farming, ecology, food processing and medicine. While sequencing devices have become ready for ubiquitous use, software tools for metagenomics still tend to be resource intensive and too hard to use for end users.

To analyze large sets of reads, classifications tools based on k -mer matching and counting such as Kraken 1 [1], Kraken 2 [2], KrakenUniq [3], Ganon [4, 5] and others are at hand. With regard to pathogen detection, KrakenUniq is particularly suited because of its very low false positive rate but also because of its efficiency and high sensitivity [3].

In order to avoid false positive classifications of reads, KrakenUniq resorts to very large databases, containing billions of encoded but uncompressed k -mers along with their respective taxa. A database is usually loaded entirely into main memory and often consumes tens of gigabytes of space. This requirement mandates specialized and expensive compute servers with very large main memory.

In recognition of this problem, KrakenUniq has been extended for database processing in chunks at read analysis time [6]. However, database creation still mandates ultralarge main memory: For example, in 2018 the production of a database for Kraken 1 and, equivalently, KrakenUniq that just covers the bacterial genomes from the RefSeq [7] “required ca. 2.5 TB of RAM and ca. 11 days” on a high-performance computer [8]. But since then, the RefSeq’s number of prokaryotic genomes alone has increased by a factor of 5.8 to over 48 Mio. [9]. Therefore, related organisms are often not covered below the species rank and also genomes below the quality level “complete genome” are usually not included in a corresponding database.

Despite the fact that these databases may not cover organisms of interest at the necessary genomic depth, users tend to apply them because otherwise, creating even larger databases would require even more main memory [8]. So, although a much richer set of genomes for organisms of interest may be available in principle – for example via recent versions of the RefSeq or in addition via Genbank [10] – these genomes cannot easily be leveraged for metagenomic analysis [11, 12].

Even though Kraken 2 offers more memory efficient k -mer databases, *it is not a suitable alternative in fields such as medical diagnostics where precision of read classification is paramount* [6]. This is because in comparison with Kraken 1 and KrakenUniq, Kraken 2 produces smaller databases, but it compresses stored k -mers under loss of information, therefore facilitating false positive k -mer counts [2, 6]. Many recent k -mer analysis tools such as Ganon follow this idea and trade k -mer compression against potentially false positive k -mer matches.

We focus on a comparison with KrakenUniq because Genestrip aims to provide a similar k -mer based functionality by avoiding k -mer compression altogether while still enabling small databases without substantial loss of precision and recall. In addition, we include Kraken 2 and Ganon (Version 2.1) in our evaluations as representatives of k -mer compressing systems.

1.2 Approach

If only a smaller group of a few dozen genera or a few 100 species is under consideration, then it seems promising to simply *generate a smaller database that contains just the k -mers from respective genomes*. However, *without additional measures, this very likely leads to many false positives*: The reason for this is that owing to the evolutionary relationship of organisms, *a large fraction of k -mers that belong to the genome of a species are unspecific* in the sense that they can as well be found in the genomes other species or genera not considered for analysis. Thus, when using too small a database, a read may be falsely assigned to a species because it could consist of k -mers that *also* belong to another species' genome not included in the database.

In order to enable a smaller database with no false positives emerging from this problem, we introduce Genestrip, as new metagenomic software tool supporting (unique) k -mer counting, read classification and read filtering.

In principle, Genestrip employs a *two-phase* database generation process: During the *entry phase*, all k -mers of genomes from selected organisms are entered into the database on the basis of the RefSeq [7] and optionally GenBank [10] as well as additional sources. During the *update phase*, all stored k -mers from the database are checked against the genomes of *all* other organisms from the RefSeq belonging to the same *category*¹. So, if a stored k -mer also occurs in another organism which is not part of the database, the k -mer's taxon still gets updated in the database and the lowest common ancestor's taxon (known as LCA) is entered instead [8]. This prevents a k -mer from being falsely associated with one of the selected organisms given that it is too unspecific.

Based on this approach, Genestrip achieves a low memory footprint during database creation as well as at read analysis time. Moreover, it offers competitive to highly favorable runtimes for k -mer matching and read classification. Together, these qualities empower users to create and use respective databases on machines with regular main memory size. Genestrip's k -mer counting results are similar to those of KrakenUniq [3] with respect to the selected organisms. Genestrip also supports unique k -mer counting with exact counts and with low memory overhead.

¹By standard, the RefSeq is structurally divided into 14 categories including "Bacteria", "Fungi", "Invertebrate", "Protozoa", "Viral" and others.

In the results section of this paper, we demonstrate the usefulness of our approach by introducing a small Genestrip database for tick-borne bacterial pathogens. An envisioned application is tick surveillance via metagenomics as reported in [13]. We performed analyses of tick-related fastq files on a regular PC in less than a minute. Thanks to the small database’s high genomic coverage for the organisms of interest, we achieved considerably better recall than with a large standard database under KrakenUniq. This means we found more k -mers and matching reads for such organisms while also sustaining high precision. As part of our performance experiments, we briefly present further small databases that may prove useful in agriculture and infectious disease diagnostics.

Genestrip is written in Java, has a small code base and supports command-line-based as well as API-based usage. Based on the Java Virtual Machine, it can be executed on modern operating systems including Linux, Windows and macOS.

Comment
on
other
K2
and
Ganon
here
too

2 Implementation

2.1 System structure

Genestrip’s database generation mandates the setting of groups of organisms whose genomes will be exclusively represented in a resulting database. Users choose the desired genera, species or strains via a configuration file by listing the corresponding taxa. If higher ranked taxa are entered, then by default, *all* of the descending taxa will be included and k -mers of their genomes will be added to the database accordingly.

Genestrip generates a database for a selected group of organisms primarily via genomes from the RefSeq [7] but genomes from Genbank and additional genomic files may be included on top or instead. By default, related genomic files are only downloaded once to a common folder and shared when further Genestrip databases are to be generated.

To better handle the complexity of the generation process, it employs a goal-oriented approach similar to Make [14]: Upon access, each goal returns a file or an object in main memory. A goal is executed only if its result is not yet available. I.e., in case of a file goal, the system first checks whether the corresponding file is present and only if it is missing, the corresponding goal is executed to produce the file.

Goal execution may involve the access of other goals which in turn, may trigger their execution. Respective goal dependencies are explicitly specified in terms of a directed acyclic graph. Details on the system’s goals and the goal graph can be found in Section 1 of the Appendix.

2.2 Database generation

Genestrip stores a database primarily as a large array L of sorted long values. Each long value encodes a k -mer from the genome of a selected organism. The default for k is 31 but smaller values and $k = 32$ are possible. In order to handle reverse complements of reads later at read analysis time, stored k -mers are standardized in the following way: Each k -mer is encoded twice as a long value — once in its canonical version and once via its reverse complement — but only the larger of the two encoded values is

stored in the database. A second array T with the same entry size as L encodes taxa for each k -mer in L at corresponding positions. At read analysis time, lookups in L are performed via a binary search.

All parts of the NCBI taxonomy tree [15] required at analysis time are also included in the database. The taxonomy extract comprises entries for the selected organisms and all of their ancestors. At analysis time, a bit vector with the same entry size as L indicates which k -mers from L were found at least once. This is sufficient for exact unique k -mer counting and is space-efficient, as it requires only about 1.6% of additional main memory with respect to L .

Genestrip employs one Bloom filter [16] at *database generation time* with a very low FPP. While the database gets filled, it tracks the already stored k -mers to avoid duplicates. Since L is not sorted at that point in time, the filter offers an efficient containment check.²

A so-called Blocked Bloom filter with a rather high FPP of almost 1% gets constructed after L got sorted. It is exclusively used for performance optimization *at read classification time*.³ The filter is consulted prior to a k -mer lookup in L to check whether the k -mer might be in L . *A negative result avoids a considerably more expensive binary search in L and therefore saves overall execution cost.* We compared the average access time to look up a random k -mer in Genestrip’s k -mer store with and without the Blocked Bloom filter being activated: Depending on the size of the database the filter brought an improvement of about factor 9 for very small databases to roughly factor 11 for larger databases. So, at read analysis time, this filter merely improves efficiency but does not cause false positives, since a hit is always backed by L . Employing the filter pays off because in practice, many k -mers from reads are likely not stored in a typical database.⁴

Regarding database generation, additional genomes can be manually added either to create deeper databases or, else, to filter out contaminating sequences that are mistakenly included in the genome of an organism of interest. For example, [19] reported that human DNA fragments often contaminate bacterial genomes for various reasons whereas [20] discussed a pipeline to remove contaminants from genomes of eukaryotic pathogens. With Genestrip, k -mers from contaminating sequences can be identified during the update phase as introduced in Section 1.2: For example, the human genome can be included via configuration in the update phase only; thus, a k -mer from the human genome that is also found in the contaminated genome will be pushed to the lowest common ancestor taxon of both species. Still, as the human genome is employed during the update phase only, a resulting database does not include k -mers from the human genome (other than those from contaminating sequences). This ensures a small-sized database.

²More recent probabilistic filters such as Xor filters [17] or a Ribbon filters [18] do not apply there since their construction demands all k -mers to be known upfront, whereas at database generation time, the k -mers stream in one by one from the genomic files being read.

³Here, we chose a Blocked Bloom filter as characterized in [17] since its query performance is superior to all other filters evaluated there. The filter’s memory overhead is tolerable considering L and T together are eight times larger in bits.

⁴For example, in the case of detecting infections, related samples typically contain a lot of host-DNA leading to many k -mers from the host in reads, whereas a related database typically would contain the pathogens k -mers only.

Of note, Genestrip allows for the exclusion of k -mers with low complexity known as genomic dust [21]: As part of the entry phase, each k -mers’s repetitiveness in terms of base pairs is measured via a fibonacci-based function. If the corresponding value exceeds a configurable threshold then the k -mer is dropped. Section 2 in the Appendix presents our dust function in detail.

2.3 Read classification

The goal `match` loads a database and performs the read classification on fastq files. The principal classification algorithm is similar to that from Kraken 1 [1] and equivalently KrakenUniq. When iterating over a read’s k -mers, each k -mer is encoded as a long – once by considering its canonical version and once via its reverse complement.⁵ For efficiency, only the larger of the two encoded values is looked up in the database. This is sufficient since, if at all, then only just the larger of two values got stored in the first place during database generation (see Section 2.2).

Genestrip implements a multithreading enabled classification procedure following a producer consumer pattern, where a producer thread reads a fastq file and puts reads into an in-memory queue while consumer threads pick reads from the queue in order to classify them.

Three types of error thresholds can be configured for read classification:

1. The ratio or number of k -mers per read that must at least be found in the database and
2. the ratio or number of k -mers per read that must be consistent with the read’s designated class taxon. A k -mer is consistent with a read’s class taxon if the k -mer’s taxon as stored in the database is the same or if it is a taxonomic ancestor of the class taxon.
3. Moreover, the classification threshold `minKMersForClass` supports the reduction of potential false positive classifications. It can be set to adjust the minimal sum of k -mers under taxon t required for a read to be classified to t . I.e., given a read r and taxon t_1 on the genus rank with two k -mers in r and taxon t_2 subordinate to t_1 on the species rank with one k -mer in r . Furthermore, r shall contain no other k -mers matching any taxons. Then, if `minKMersForClass=2`, r would *not* be classified to t_1 but to t_2 instead since the single k -mer under t_1 is below the threshold but the sum of three k -mers under t_2 is sufficient.

The thresholds allow for playing off recall against precision in light of potentially high error rates of some sequencing devices [22, 23]. For all the experiments throughout this paper, the first two thresholds were left at the default of one k -mer which is the most sensitive level.

2.4 Analysis reports

The read classification via the goal `match` produces analysis reports in CSV file format with a comprehensive set of statistical values. We use the following notation to describe

⁵Considering both versions is necessary because sequenced reads may occur in either direction but also because genomes may be represented in reverse complement direction, such as in the case of circular DNA.

them: Let $k_{t,f,d}$ be the number of k -mers for taxon t found in fastq file f with respect to database d and let $u_{t,f,d}$ be the corresponding number of unique k -mers. Similarly, let $u_{t,d}$ be the number of (unique) k -mers for t stored in d .

Table 1 describes the most important reported columns. Additional columns exist to accumulate k -mer counts and read counts etc. upward to ancestors in the taxonomy. We also track the standard deviations of `avg. contig length`, `avg. read length`, `mean class error` and other columns wherever suitable. Moreover, normalized values derived from the original values by dividing by $u_{t,d}$ are included where suitable. Normalized values aim to reduce counting bias due to a high or low representation of a taxon t in terms of its stored k -mers $u_{t,d}$.

Column name	Description
<code>taxid</code>	The class taxon t
<code>reads</code>	The number of reads classified with respect to t
<code>kmers from reads</code>	The number of k -mers from reads for t that are consistent with t
<code>kmers</code>	$k_{t,f,d}$
<code>unique kmers</code>	$u_{t,d}$
<code>avg. contig length</code>	The average length of contiguous sequences of k -mers that are specific to t 's genome(s) as represented in d
<code>max contig length</code>	The maximum length of all contiguous sequences of k -mers that are specific to t 's genome(s)
<code>avg. read length</code>	The average length of reads for t in base pairs
<code>db coverage</code>	$u_{t,f}/u_{t,d}$
<code>mean class error</code>	The mean ratio of k -mers from a read r for t , where those k -mers are not consistent with t per r 's total number of k -mers

Table 1 Extract of columns from Genestrip's read analysis reports

2.5 Example Usage

We give a brief impression on how to use Genestrip from the Unix command line.⁶ The installation includes a sample project `human_virus`. After building Genestrip, you may call

```
sh ./bin/genestrip.sh human_virus dbinfo
```

from the installation directory to generate the `human_virus` database and also, to create a CSV file with basic information about the database content.

As stated in Section 2.1, Genestrip follows a goal-oriented approach in order to create any result files. So, when generating the `human_virus` database, Genestrip will automatically

1. download the NCBI taxonomy and unzip it to `./data/common`, if not yet present,
2. download the RefSeq release catalog to `./data/common/refseq`, if not yet present,
3. download all virus related RefSeq fna files to `./data/commmon/refseq` (which is currently just a single g-zipped file), if not yet present,

⁶For detailed information see: <https://github.com/pfeiferd/genestrip/blob/master/README.md>

4. perform several follow-up goals, until the database file `human_virus_db.zip` is finally made under `./data/projects/human_virus/db` and
5. create a CSV file `human_virus_dbinfo.csv` under `./data/projects/human_virus/csv`, which contains basic information about the database, i.e. the number of k -mers stored per taxon.

The generated database comprises k -mers for all viruses whose taxons are listed in the configuration file `./data/project/human_virus/taxids.txt`.

To try out classification, the installation comes with the small fastq file `sample.fastq.gz` in `./data/projects/human_virus/fastq`. One may trigger its analysis via

```
sh ./bin/genestrip.sh human_virus match -f ./data/projects/
↪ human_virus/fastq/sample.fastq.gz
```

The resulting CSV file named `human_virus_match_sample.csv` can then be found under `./data/projects/human_virus/csv`.

3 Results

3.1 Correctness

To assess Genestrip’s correctness with regard to k -mer counting and read classification, we compared example databases and respective analysis results directly and in detail with those from KrakenUniq. As part of our test procedure, we created database contents under both systems that were identical in terms of their genomic basis and their stored k -mers. We then applied the databases to suitable fastq files in order to compare the results between both systems.⁷ Highly similar results from this check mutually solidify the correctness of both systems since they had been developed completely independently. (In addition, Genestrip’s correctness is also supported via extensive automated functional tests with high test coverage as part of its code base.)

We chose all RefSeq genomes from the category “Viral” for the two databases. The latter turned out to be completely identical in terms of the number of k -mers stored per taxon. In the following, we name the two databases `cv_ku` and `cv_g` with regard to KrakenUniq and Genestrip, respectively.

We applied `cv_ku` and `cv_g` to a simulated fastq file extracted from [24] where the file’s reads were used to benchmark metagenomic pipelines aimed at the discovery of new viruses. The file contains 61,550 long reads extracted from 6,155 viral standard genomes from the RefSeq [24]. Furthermore, we generated simulated fastq files based on all RefSeq genomes from the category “Viral” via the tool InSilicoSeq [25]: We applied the respective Illumina generation models “MiSeq” and “HiSeq” to produce two pairs of fastq files with one million reads per pair. Then, we analyzed these files via `cv_ku` and `cv_g` as well. In all of the cases the k -mer counts per taxon were identical between KrakenUniq and Genestrip and the classified read counts were also identical.

As real-world examples, we applied the two databases to short-read fastq files from DNA-sequencing human saliva of five different donors. Table 2 shows basic

⁷For compatibility with KrakenUniq, we kept `minKMersForClass` from Section 2.3 at the default of 1.

statistics of these files as downloaded from the Sequence Read Archive [26]. Again, we achieved completely identical k -mer and read counts between KrakenUniq and Genestrip. Details on some of these comparisons can be found in Section 3 of the Appendix.

#	Base file name	Reads	Avg. BPs per read
1	SRR5571985	812,085,208	101
2	ERR1395613	900,709,176	191
3	SRR5571991	980,879,835	276
4	ERR1395610	824,479,570	429
5	SRR5571990	605,561,636	685

Table 2 Real-world fastq files from DNA-sequencing five human saliva samples as used throughout the evaluation.

As previously stated, this perfect agreement mutually solidifies the correctness of both systems, but obviously, it does not yet highlight the advantages of Genestrip when creating and using small databases for selected groups of organisms.

3.2 Classification via broad coverage databases

As a baseline, we tested read classification quality using all of the simulated data introduced in Section 3.1. With the ground truth known, this allowed for comparing precision and recall across different k -mer-based analysis tools.

Regarding the databases, we reused `cv_ku` and `cv_g` from Section 3.1 and generated analogous databases for Kraken 2 and Ganon, respectively. Ganon depends on an important hyper-parameter `max-fp` that controls the false positive probability (FPP) for matching k -mers and must be set at database generation time. (A lowered FPP incurs a larger database.) We used the default value of 0.001 (combined with `--filter-type hibf`) which is documented as highly sensitive in Ganon’s online manual.⁸

All analysis tools but KrakenUniq offer classification time options to trade sensitivity against confidence. As listed in Table 3, we examined two variants for each tool – a variant with higher sensitivity indicated by “(S)” (which is the default setting for Kraken 2 and Genestrip) and a variant with higher confidence indicated by “(C)”. Ganon performs multiple classifications for reads – typically in less than 15% of the cases. To enable a reasonably fair comparison, we computed the lowest common ancestor when Ganon produced multiple class taxons for a read and used the outcome in the evaluations.

Table 5 shows that most systems perform within reason except for the configuration “Kraken 2 (C)”. Irrespective of its configuration, Ganon appears to be sensitive to the higher error rates in “MiSeq” as opposed to “HiSeq”.

⁸See https://pivovc.github.io/ganon/default_databases/#filter-type-ibf-and-hibf

System variant	Description	Option settings
KrakenUniq	High sensitivity (default)	none
Kraken 2 (S)	High sensitivity (default)	none
Kraken 2 (C)	High confidence	--confidence 0.5
Ganon (S)	High sensitivity	--rel-cutoff 0 --rel-filter 0
Ganon (C)	High confidence (default)	none
Genestrip (S)	High sensitivity (default)	none
Genestrip (C)	High confidence	minKMersForClass=4

Table 3 Classification time settings for each tool: A higher sensitivity and a higher confidence variant are examined in most of the experiments from Section 3.

KrakenUniq does not offer related options and therefore, its variant name carries no “(S)” or “(C)”.

System	Data	Species			Genus		
		Prec.	Recall	F ₁	Prec.	Recall	F ₁
KrakenUniq	Extr. from [24]	0.95	0.95	0.95	1.00	1.00	1.00
Kraken 2 (S)	Extr. from [24]	0.95	0.95	0.95	1.00	1.00	1.00
Kraken 2 (C)	Extr. from [24]	0.71	0.71	0.71	0.95	0.95	0.95
Ganon (C)	Extr. from [24]	0.96	0.96	0.96	1.00	1.00	1.00
Ganon (S)	Extr. from [24]	0.97	0.97	0.97	1.00	1.00	1.00
Genestrip (S)	Extr. from [24]	0.95	0.95	0.95	1.00	1.00	1.00
Genestrip (C)	Extr. from [24]	0.95	0.95	0.95	1.00	1.00	1.00
KrakenUniq	MiSeq	0.91	0.91	0.91	0.99	0.99	0.99
Kraken 2 (S)	MiSeq	0.91	0.91	0.91	0.99	0.99	0.99
Kraken 2 (C)	MiSeq	0.77	0.76	0.76	0.96	0.93	0.94
Ganon (C)	MiSeq	0.92	0.77	0.83	0.99	0.83	0.90
Ganon (S)	MiSeq	0.92	0.92	0.92	0.99	0.99	0.99
Genestrip (S)	MiSeq	0.91	0.91	0.91	0.99	0.99	0.99
Genestrip (C)	MiSeq	0.91	0.91	0.91	0.99	0.99	0.99
KrakenUniq	HiSeq	0.88	0.88	0.88	0.98	0.98	0.98
Kraken 2 (S)	HiSeq	0.88	0.88	0.88	0.98	0.98	0.98
Kraken 2 (C)	HiSeq	0.79	0.78	0.79	0.96	0.95	0.95
Ganon (C)	HiSeq	0.88	0.83	0.86	0.98	0.93	0.96
Ganon (S)	HiSeq	0.89	0.89	0.89	0.98	0.98	0.98
Genestrip (S)	HiSeq	0.88	0.88	0.88	0.98	0.98	0.98
Genestrip (C)	HiSeq	0.88	0.88	0.88	0.98	0.98	0.98

Table 4 Read classification quality for different simulated fastq files analyzed with KrakenUniq, Kraken 2, Ganon and Genestrip. Regarding KrakenUniq and Genestrip the viral databases `cv_ku` and `cv_g` from Section 3.1 were used. Similar (unnamed) viral databases were generated for Kraken 2 und Ganon. The systems’ adaption for “(S)” and “(C)” are based on Table 3.

3.3 Classification via small databases

The following results highlight the importance of the update phase according to Section 1.2 in the context of small databases.

3.3.1 Human virus vs. complete viral databases

As a first example, we created two small viral databases based on the RefSeq covering *human viruses only*⁹, again under both KrakenUniq and Genestrip. In the following, we name the two databases `hv_ku` and `hv_g` with regard to KrakenUniq and Genestrip, respectively. We again assured that the databases were identical in terms of their genomic basis and their stored k -mers. In addition, we generated analogous human virus databases for Kraken 2 and Ganon with the same configurations for Ganon as described in Section 3.2.

Figure 1 compares the viral database contents for the databases `cv_ku` and `cv_g` from Section 3.1 against the respective “human virus only” database for (a) KrakenUniq and (b) Genestrip. In addition, (c) directly compares the k -mer counts of the two “human virus only” databases `hv_ku` and `hv_g`. The plots each depict data points of the 422 taxa referenced in at least one of the two corresponding databases.

Regarding Genestrip’s databases from Figure 1 (b), nearly all k -mer counts on the species rank and below are identical per taxon as an effect of the update phase (that we discussed in Section 1). This is an important precondition met by Genestrip’s `hv_g` database to avoid false positives at read classification time. Nevertheless, data points from higher taxonomic ranks scatter because the complete viral database `cv_g` contains k -mers from many more non-human viruses, and these k -mers were pushed to higher ranks due to the lowest common ancestor update. Obviously, these additional k -mers are not required to classify reads accurately with respect to human viruses. The database `hv_g` contains only 3,865,987 k -mers with a disk size of approximately 32 MB, whereas `cv_g` has a disk size of 3.1 GB with 98 times as many k -mers.

KrakenUniq’s corresponding k -mer counts from Figure 1 (a) are not identical for the species rank and below. Although visually not well obvious in (a), this accounts for more than 150 scattering data points on the species rank and below. These differences offer the potential for more error and in particular false positives at read classification time as will be exemplified next. The plots (a) and (c) look similar because with regard to human virus taxa, `hv_g` is similar to `cv_ku` and `cv_g`.

3.3.2 Small databases on simulated data

We applied the “human virus only” databases to the “MiSeq” and “HiSeq” data from Section 3.1 and again measured precision and recall but with respect to human viruses only.¹⁰ So, this evaluation examines the quality of a small database on simulated data where the ground truth is known but with a focus on just the genomes and taxons used to create the small databases. Since about 98% of the reads in “MiSeq” and “HiSeq” data are from non-human-related viruses, they constitute noise in this context.

Table 5 shows that all systems still perform within reason on the simulated data but with Genestrip clearly in the lead.

⁹The list of human viruses originates from [27]. Some human viruses known to be missing were manually added. All taxons were chosen on the species rank.

¹⁰In this context, we do not present the results for the extracted fastq file from [24] since the file contains only 240 reads from human viruses which caused low precision outliers under KrakenUniq and Kraken 2. We therefore omit the file also below when focusing the evaluation on human viruses.




Fig. 1 Scatter plots of *k*-mers per taxon in a complete viral database versus a “human virus only” database both from the RefSeq created via (a) KrakenUniq and (b) Genestrip (based on a logarithmic scale with 1 added to each value to enable the representation of zero-counts). In addition, (c) directly compares the *k*-mer counts of the two “human virus only” databases *hv_ku* and *hv_g*.

3.3.3 Small databases on real-world data

As real-world examples, we first applied the “human virus only” databases to the short-read fastq files from Table 2 (and introduced in Section 3.1).

Figure 2 (a,b) shows the scatter plots of *k*-mer counts and read counts from KrakenUniq on a logarithmic scale for the database *cv_ku* versus *hv_ku* when applied to the human saliva fastq file #3 from Table 2. Figure 2 (c,d) shows the corresponding results with respect to the Genestrip databases *cv_g* versus *hv_g*: While the small database under KrakenUniq led to massive overcounts of *k*-mers for the rank species

System	Data	Species			Genus		
		Prec.	Recall	F ₁	Prec.	Recall	F ₁
KrakenUniq	MiSeq	0.71	0.96	0.82	0.71	0.96	0.82
Kraken 2 (S)	MiSeq	0.70	0.96	0.81	0.70	0.96	0.81
Kraken 2 (C)	MiSeq	0.95	0.91	0.93	0.95	0.91	0.93
Ganon (C)	MiSeq	0.95	0.80	0.87	0.95	0.80	0.87
Ganon (S)	MiSeq	0.37	0.96	0.53	0.37	0.96	0.53
Genestrip (S)	MiSeq	0.95	0.96	0.95	0.95	0.96	0.96
Genestrip (C)	MiSeq	0.95	0.96	0.96	0.96	0.96	0.96
KrakenUniq	HiSeq	0.82	0.97	0.89	0.82	0.97	0.89
Kraken 2 (S)	HiSeq	0.82	0.97	0.89	0.82	0.97	0.89
Kraken 2 (C)	HiSeq	0.96	0.95	0.95	0.96	0.95	0.95
Ganon (C)	HiSeq	0.96	0.93	0.94	0.96	0.93	0.95
Ganon (S)	HiSeq	0.68	0.97	0.80	0.68	0.98	0.81
Genestrip (S)	HiSeq	0.97	0.97	0.97	0.98	0.98	0.98
Genestrip (C)	HiSeq	0.97	0.97	0.97	0.98	0.98	0.98

Table 5 Read classification quality for simulated fastq files from Section 3.1 analyzed with KrakenUniq, Kraken 2, Ganon and Genestrip using respective “human virus only” databases and with an evaluation focus on human viruses only. The systems’ adaptations for “(S)” and “(C)” are based on Table 3.

and below, this did not happen under Genestrip. Moreover, KrakenUniq’s overcounting of k -mers strongly influenced the counts of classified reads. Regarding Genestrip, the read counts on the species level and below are affected only by differing k -mer counts that occurred on the genus level and above. However, the impact remains small. The results for the other human saliva fastq files are similar (not shown).

We also measured a relative read classification quality for all of the five human saliva fastq files from Table 2 well as for simulated fastq files from Section 3.1 by comparing the “human virus only” database against the corresponding complete viral database and under KrakenUniq, Kraken 2, Ganon and Genestrip – again with an evaluation focus on human viruses only. Although no ground truth is available for the human saliva fastq files, this comparison tells *how much classification quality suffers under each analysis tool when the corresponding database is filled with just the genomes of interest – i.e., the genomes of human viruses only instead of all viral genomes*. The result represents a relative factorial change where an approximate baseline for the simulated data is given by Table 4 and by Table 3 of the Appendix.

Table 6 shows the results for simulated data and fastq file #3 from Table 2. The results for the other four human saliva fastq files are similar and supplied in Table 4 of the Appendix.

While KrakenUniq and Ganon still achieve a mediocre to good F_1 -score on simulated data, their precision is insufficient regarding the human saliva fastq files. This is due to massive false positive classifications of reads, and is a consequence of a missing consideration of non-human-related viral genomes during database generation. Although speculative, we assume that the real-world data has a different error profile than the simulated data, which may contribute to the stark difference. Also, reads belonging to host DNA may cause critical noise. “Kraken 2 (C)” offers reasonable results, but “Genestrip (C)” starts from a better baseline in Table 4, and its human




Fig. 2 Scatter plots of k -mer counts (a, c) and read counts (b, d) for the human saliva fastq file #3 from Table 2 analyzed with KrakenUniq (a, b) and Genestrip (c, d) using the complete viral databases versus the “human virus only” databases, as both shown in Figure 1 (based on a logarithmic scale with 1 added to each value to enable the representation of zero-counts).

virus database holds up because it does consider non-human-related viral genomes as part of the update phase according to Section 1.2.

3.4 Example: Tick-borne bacterial pathogens

This section presents the detection of tick-related bacterial pathogens as an example application of Genestrip. Once again, we compare our tool against KrakenUniq,

System	Data	Species			Genus		
		Prec.	Recall	F ₁	Prec.	Recall	F ₁
KrakenUniq	MiSeq	0.71	0.96	0.82	0.71	0.96	0.82
Kraken 2 (S)	MiSeq	0.70	0.96	0.81	0.70	0.96	0.81
Kraken 2 (C)	MiSeq	0.93	0.95	0.94	0.93	0.95	0.94
Ganon (C)	MiSeq	0.95	0.96	0.95	0.95	0.96	0.95
Ganon (S)	MiSeq	0.37	0.96	0.53	0.37	0.96	0.53
Genestrip (S)	MiSeq	0.95	0.96	0.96	0.95	0.96	0.96
Genestrip (C)	MiSeq	0.96	0.96	0.96	0.96	0.96	0.96
KrakenUniq	HiSeq	0.82	0.98	0.89	0.82	0.98	0.89
Kraken 2 (S)	HiSeq	0.81	0.98	0.89	0.81	0.98	0.89
Kraken 2 (C)	HiSeq	0.94	0.98	0.96	0.94	0.98	0.96
Ganon (C)	HiSeq	0.96	0.98	0.97	0.96	0.98	0.97
Ganon (S)	HiSeq	0.68	0.98	0.80	0.68	0.98	0.80
Genestrip (S)	HiSeq	0.98	0.98	0.98	0.98	0.98	0.98
Genestrip (C)	HiSeq	0.98	0.98	0.98	0.98	0.98	0.98
KrakenUniq	#3 from Table 2	0.04	0.91	0.07	0.04	0.91	0.07
Kraken 2 (S)	#3 from Table 2	0.03	0.65	0.06	0.03	0.65	0.06
Kraken 2 (C)	#3 from Table 2	0.40	1.00	0.57	0.40	1.00	0.57
Ganon (C)	#3 from Table 2	0.25	1.00	0.40	0.25	1.00	0.40
Ganon (S)	#3 from Table 2	0.02	0.69	0.04	0.02	0.70	0.04
Genestrip (S)	#3 from Table 2	0.50	1.00	0.67	0.50	1.00	0.67
Genestrip (C)	#3 from Table 2	0.93	1.00	0.96	0.93	1.00	0.96

Table 6 Read classification quality for simulated fastq files from Section 3.1 as well as the real-world fastq file #3 from Table 2 analyzed with KrakenUniq, Kraken 2, Ganon and Genestrip with an evaluation focus on human viruses only. The evaluation compares the classifications when using a “human virus only” database against each system’s own classifications obtained when using the corresponding complete viral database from Section 3.2. The systems’ adaptations for “(S)” and “(C)” are based on Table 3.

Kraken 2 and Ganon reuse or generate standard databases available for the latter systems. Regarding Genestrip, we employ a relatively small database **tb** with high genomic coverage for the organisms of interest.

An envisioned application is tick surveillance via metagenomics according to [13], where the authors applied Nanopore technology to check eight field-collected ticks but back then relied on Kraken 2 to identify corresponding organisms. Due to [28], [29] and [30] we also see a use case when screening for tick-borne pathogens in vertebrates potentially including humans.

3.4.1 Analysis databases

We discuss the databases used in conjunction with each system but with an emphasis on Genestrip’s database **tb**.

MicrobialDB [31] for KrakenUniq aims at the diagnosis of infectious diseases using metagenomics and contains about 98 times as many k -mers as **tb**. It has a disk size of 574 GB and covers a broad set of microbials including bacteria and the human genome with contaminants removed [20]. Regarding Kraken 2, we used the latest **Standard** database from [31] covering the entire RefSeq category “Bacteria”, the human genome and more with a disk size of 73 GB and with about 18 billion compressed k -mers.

For Ganon, we generated a database covering all the genomes of category “Bacteria” from the RefSeq Release 233 using Ganon’s default FPP (see also Section 3.2) which resulted in a database size of 58 GB.

The database **tb** for Genestrip is based on the RefSeq Release 233 and covers all RefSeq genomes under the twelve bacterial genera *Bartonella*, *Rickettsia*, *Anaplasma*, *Ehrlichia*, *Borrelia*, *Borrelia*, *Chlamydia*, *Mycoplasma*, *Malacoplasma*, *Mycoplasma*, *Metamycoplasma* and *Francisella* as many included species are pathogenic to humans and known or assumed to occur in ticks [32–36]. To avoid contaminants from human DNA and black-legged tick DNA, we added corresponding genomes to Genestrip’s update phase as explained at the end of Section 2.2. We also activated the genomic dust filter at moderate sensitivity – e.g., the k -mer `TTTTTTTTTTTTTCGGTGTGTTTTGTTACT` would have been dropped because its repetitiveness is just above the set threshold.

tb contains about 525 million uncompressed k -mers and has a disk size of just 3.91 GB including the accompanying Bloom filter and a suitable extract of the NCBI taxonomy tree. **tb** has 1.96 times as many k -mers stored under the selected genera and below as *MicrobialDB*. The main reason behind this increase is the general growth of the RefSeq’s number of prokaryotic genomes between the generation of *MicrobialDB* in mid 2023 and the generation of **tb** [9].

The generation of **tb** revealed that its included genomes from the RefSeq have a high overlap: On average each stored k -mer occurred 11 times in these genomes. Only 0.0079 % of all k -mers considered for storing were dropped as “dust”. The update phase pushed 26.3 % of the stored k -mers to ranks above genus. 0.027 % of the stored k -mers had been pushed to the taxon of “cellular organisms” because tick DNA and human DNA were included on purpose during the update phase for the sake for contaminant detection.

3.4.2 Analysis of simulated data

To assess classification quality for this case, we first generated ground truth data via the eight real-world fastq files from [13] mentioned above. For this purpose, we employed the tool NanoSim from [37, 38] which is capable of generating realistic long reads based on a set genomes and a sample fastq file used for training. The tool tries to reflect the genomic abundance of organisms from the training fastq file. Moreover, it imitates long read sequencing errors resulting from Nanopore technology via an error model which is also estimated on the basis of the training fastq file.

We supplied NanoSim with all the RefSeq genomes from the twelve bacterial genera mentioned in Section 3.4.1 and used each of the eight tick-related real-world fastq files to generate corresponding *simulated* fastq files. The number of reads per simulated file ranges between 13,000 and 160,000 with 319,531 reads in total and an average read length of 4,480 BPs. It is important to know that both, the real-world data and the simulated data are subject to high error in their bases: As determined with Genestrip itself, the mean ratio of k -mers that could not be assigned to any taxon in classified reads was 83 % and 88 %, respectively. Regarding the simulated data, the error cannot be attributed to missing k -mers in **tb**, since the simulated data was generated only via genomes also represented in **tb**.

We applied the analysis tools with their respective databases to the eight simulated fastq files. Apparently due to the high error rates in the data, Ganon did not classify any reads under the variant “Ganon (C)” – hence, we show no results for this case. Instead, we present the additional variant “Ganon (A)” where thresholding is completely turned off via the combined options `-rel-cutoff 0, --rel-filter 0` and `--fpr-query 1e-0` causing all reads to be classified.

Table 7 shows the marco-averaged results for each system. Apparently, they all struggle with the high error from the modeled Nanopore device but Genestrip’s configurations perform far better than the others. We attribute this to the deep genomic coverage of the organisms of interest in `tb` combined with a loss-less storing of related k -mers and Genestrip’s thorough least common ancestor update.

System	Species			Genus		
	Prec.	Recall	F ₁	Prec.	Recall	F ₁
KrakenUniq	0.44	0.15	0.21	1.00	0.26	0.39
Kraken 2 (S)	0.53	0.19	0.27	1.00	0.32	0.47
Kraken 2 (C)	0.33	0.00	0.01	1.00	0.02	0.03
Ganon (S)	0.35	0.18	0.23	0.62	0.29	0.38
Ganon (A)	0.18	0.18	0.18	0.29	0.29	0.29
Genestrip (S)	0.60	0.44	0.51	0.79	0.59	0.67
Genestrip (C)	0.58	0.40	0.47	0.77	0.54	0.63

Table 7 Macro averaged read classification quality for simulated data derived from analysis of bacterial pathogens in ticks [13]. The results are based on eight simulated fastq files generated using NanoSim [37, 38] on the basis of eight corresponding real-world, tick-related fastq files. Section 3.4.1 describes the databases used and Section 3.4.2 discusses the analyzed data. The systems’ adaptations for “(S)” and “(C)” are based on Table 3; “Ganon (A)” is described in Section 3.4.2.

3.4.3 Explorative analysis of tick-related data

In order to investigate the effect of databases with deep genomic coverage, we applied KrakenUniq and Genestrip along with MicrobialDB and `tb` to the eight *real-world* fastq files from [13] mentioned above. For this comparison we take advantage of the compatibility of Genestrip and KrakenUniq as established in Section 3.1.¹¹ Then, differences of classification results solely originate in the databases – but as explained in Section 3.4.1, `tb` is deeper than MicrobialDB as it has almost twice as many k -mers stored for its included genera.

As an example result, Figure 3 shows the five most frequent k -mer counts and reads counts for “Tick 6” accumulated up to the genus rank. In all the cases but “138” (*Borrelia*), the aggregated counts for the genera under `tb` are mildly to considerably higher than those achieved with MicrobialDB. We observed similar improvements in terms of k -mers and reads with respect to *Ehrlichia*, *Rickettsia* and partly to *Anaplasma* in

¹¹To ensure compatibility for this experiment, we set `minKMersForClass=1` under Genestrip (see Section 2.3).

all of the eight ticks’ fastq files whenever the corresponding organisms showed presence via **MicrobialDB**. The effect is consistent with Genestrip’s higher recall on the simulated data from Table 7.




Fig. 3 Comparison of *k*-mer counts (a) and read counts (b) between the database **MicrobialDB** (based on KrakenUniq) and **tb** (based on Genestrip). The counts are aggregated up to rank genus for the long read fastq file of “Tick 6” from [13] and restricted to the genera included in **tb** with only the five most frequent genera in terms of *k*-mer counts shown.

Figure 4 details the results for Tick 6 with respect to “943” (*Ehrlichia*): Interestingly, Genestrip assigns most of the found *k*-mers right to the genus, whereas KrakenUniq produces a peak under the strain “1423892” (*Ehrlichia muris* AS145). The difference originates in the deeper database **tb**, where in comparison with **MicrobialDB**, more common *k*-mers were found from a broader range of genomic data under the genus *Ehrlichia* at **tb**’s generation time. In turn, Genestrip’s update phase pushed these *k*-mers to the genus rank as the lowest common ancestor. Indeed, the ratio of stored *k*-mers between **tb** and **MicrobialDB** is 4.54 for the taxon “943” but only 0.24 for “1423892”.

Again, we observed similar results, particularly with respect to *Ehrlichia*, *Rickettsia* and partly to *Anaplasma*, in all of the eight ticks’ fastq files whenever corresponding organisms were present via **MicrobialDB**. These findings form a general trend when *k*-mer databases with a lowest common ancestor representation become deeper. The effect was first discovered and studied in [8]. Note that this is beneficial since such deeper *k*-mer databases tend to protect read classification from potentially falsely assigning a too low-ranked taxon.

3.5 Performance

As mentioned in Section 1, Genestrip’s focus is on generating and using potentially small but also deep *k*-mer databases for metagenomic analysis of selected organisms of interest. This implies that users must be empowered to flexibly create new databases – ideally on their own devices. We therefore verify that Genestrip’s database generation




Fig. 4 Comparison of k -mer counts (a) and read counts (b) between the database **MicrobialDB** (based on KrakenUniq) and **tb** (based on Genestrip) for all taxons under the genus “943” (*Ehrlichia*) including the genus itself. The counts belong to the long read fastq file of “Tick 6” from [13] – the same file as in Figure 3.

performance adheres to this requirement. Moreover, we report on read classification performance in comparison with KrakenUniq, Kraken 2 and Ganon.

The following experiments were all performed on a modern PC with 128 GB of RAM, 10 CPU cores and a 2 TB solid state disk. Although for Genestrip about 32 GB of RAM would have sufficed, we needed the additional RAM for the competing analysis tools.¹²

3.5.1 Database generation

Table 8 lists six Genestrip databases used for assessing database generation performance. They refer to several categories and were designed with a potential use in medicine or agriculture in mind. Table 9 reports execution statistics for the generation of these databases from scratch. The execution times and the maximum RAM remained moderate in all of the cases.¹³

The different execution times in Table 9 depend on the size of the database itself but also on the RefSeq category, where “Bacteria” is the largest of the five ones used with genomes in over 3,250 zip files with a total disk size of over 575 GB (in RefSeq Release 233). The zip files must all be screened during each goal from Table 1 of the Appendix but the update phase, i.e., the goal `updatedb`, dominates the total execution time.

¹²We ensured that as few competing processes as possible were running during the experiments and OS disk swapping was disabled. Platform details: Intel Xeon W-2255 CPU @3.70 GHz, Ubuntu 24.04.4 LTS, Java Oracle 23.0.1 64-Bit Server VM.

¹³Times for downloading any of the necessary base files from the NCBI are not included. RefSeq and Genbank files were downloaded upfront via the goals `refseqfna` and `fastasgenbankdl` from Figure 1 of the Appendix. The databases were generated via the goal `db` from Figure 1 of the Appendix.

DB name	Category	Description	Species and below
cv_g	Viral	See Section 3.1	25,739
hv_g	Viral	See Section 3.3	225
tb	Bacterial, Plasmid	See Section 3.4.1	1,119
protozoa	Protozoa	Comprehensive collection of protozoans from the RefSeq which are pathogenic to humans [39]	108
vineyard	Fungi	Fungal infections from the RefSeq and Genbank threatening grapevine plants [40]	11
parasites	Invertebrate	Comprehensive collection of parasitic invertebrates from the RefSeq which are pathogenic to humans [39]	77

Table 8 Background on the Genestrip databases used for assessing database generation performance.

DB name	Stored k -mers	Disk size (MB)	Wall time (min)	Max. RAM (MB)
cv_g	379,165,091	3,218	47	7,918
hv_g	3,865,987	32	5	1,114
tb	524,530,554	3,731	489	23,003
protozoa	1,332,774,884	9,621	138	25,306
vineyard	688,413,684	4,662	87	24,681
parasites	1,659,904,652	11,509	278	30,636

Table 9 Statistics for generating the six Genestrip databases from Table 8 using 20 threads and a maximum allowed JVM heap size of 32 GB.

3.5.2 Read classification

To assess Genestrip’s performance regarding k -mer counting and read classification, we compared its analysis runs with corresponding runs under KrakenUniq, Kraken 2 and Ganon.

For the first comparison we applied the viral databases from Section 3.2 to the human saliva fastq files from Table 2. Scenario (a) of Table 10 summarizes the analysis performance across all five fastq files.

In scenario (b) of Table 10 we used the bacterial databases like **MicrobialDB** and **Standard** from Section 3.4.1 vs. **tb** to analyze the eight tick-related fastq files from Section 3.4.3. On the one hand, comparing the tools this way may be considered “unfair” because of unequally sized databases. But for KrakenUniq, Kraken 2 and Ganon, there is no good alternative to employing broad-coverage bacterial databases for the purpose (of detecting tick-related bacterial pathogens) because, as examined in Section 3.3, a corresponding small database bears the high risk of many false positives. Since **MicrobialDB** is so large, we had to make use of KrakenUniq’s database processing in chunks which broke down each run into 6 rounds. This caused considerably longer run times under KrakenUniq.

In scenario (c) we had the fastq files from (b) jointly analyzed. In this case,, Kraken 2 and Genestrip both benefitted considerably from loading the database just once for all eight fastq file, whereas Ganon profited only mildly and KrakenUniq not at all.

Genestrip dominates across all three scenarios with respect to classification speed.

Regarding main memory, all the tools’ analysis performance is dominated by the size of the database. The required memory did not vary much across runs when the same database was used. Scenario (a) fairly compares the memory footprint of each tool since all databases were generated via the same set of genomes: As expected Kraken 2 and Ganon need less space than KrakenUniq and Genestrip because of their compressed representation of k -mers. Interestingly, Genestrip needs much less space than KrakenUniq still, although they both work with a similar, uncompressed representation of k -mers.

Scenario	Parameter	KrakenUniq	Kraken 2	Ganon	Genestrip
(a) Viral DBs from Section 3.2 on the 5 human saliva fastq files from Table 2	Avg. reads per fastq file		824,743,085		
	Avg. BPs per read		101		
	Avg. RAM (MB)	10,617	1,410	2,472	4,560
	Max. wall time (min)	38	40	147	37
	Avg. wall time (min)	31	34	124	30
	Avg. speed (reads / s)	442,221	402,313	110,852	458,190
	Min. reads per fastq file		644,767		
(b) broad-coverage bacterial DBs from Section 3.4.1 vs. <code>tb</code> on the 8 tick-related fastq files from Section 3.4.3	Max. reads per fastq file		2,233,967		
	Avg. reads per fastq file		1,255,973		
	Avg. BPs per read		365		
	Avg. RAM (MB)	121,676	96,866	58,380	6,148
	Max. wall time (min)	25.4	3.35	2.68	0.81
	Avg. wall time (min)	22.1	3.22	1.83	0.76
	Avg. speed (reads / s)	946	6,508	11,523	27,303
(c) broad-coverage bacterial DBs from Section 3.4.1 vs. <code>tb</code> on the 8 tick-related fastq files <i>jointly analyzed</i>	RAM (MB)	121,775	97,246	58,410	6,189
	Wall time (min)	180	4.8	9.61	1.5
	Speed (reads / s)	928	34,888	17,413	116,642

Table 10 Read analysis performance comparison between KrakenUniq, Kraken 2, Ganon and Genestrip with 10 threads for each tool. For MicrobialDB we set the parameter `--preload-size 96G` to limit KrakenUniq’s main memory according to [6]. (Otherwise, the analysis with KrakenUniq would have not been possible on the given PC.)

4 Conclusions

We have presented Genestrip – a bioinformatics tool for (unique) k -mer counting and read classification. Its major improvement over existing systems is its very low memory consumption at database generation time and at read analysis time while also maintaining very high precision and recall. To achieve this, Genestrip facilitates and suggests the generation of small databases that cover the genomes of a collection of organisms of interest. Genestrip’s update phase ensures that during database generation, the most suitable lowest common ancestor taxon is computed for each stored k -mer by *also considering the genomes of organisms whose k -mers are not included in the database*. Normally, *Genestrip considers all genomes from a matching category of the RefSeq for this purpose*.

The results from Section 3.1 indicate that Genestrip attains the same type of correctness as KrakenUniq in terms of database structure, k -mer counting and read classification given that the databases of both systems are filled via all genomes from an entire RefSeq category. But more importantly, Section 3.3 highlights that Genestrip sustains high classification quality when changing from a large, comprehensive database to a small database, although just a subset of the genomes out of a RefSeq category is included in that small database. In contrast, we showed in Section 3.3.3 that corresponding small databases under KrakenUniq, Kraken 2 or Ganon fail to achieve a suitable level of read classification quality. This is because the latter tools do not consider genomes outside of the database to be generated like Genestrip does during its update phase at database generation time.

Section 3.4 has exemplified the usefulness of Genestrip’s “small databases” in the context of tick surveillance where ticks were screened for a collection of well-known bacterial pathogens via metagenomics [13]. The resulting database `tb` has a disk size of less than 4 GB and can be used for analysis on any modern PC. Nevertheless, Genestrip offers the exactness of KrakenUniq’s very large database `MicrobialDB` in terms of k -mer counting, unique k -mer analysis and read classification for the chosen genera.

Unlike Kraken 2 and Ganon, Genestrip also eliminates the potential for false positive k -mer counts because it does not compress k -mers under loss of information. This renders our tool suitable in fields such as medical diagnostics where precision of k -mer counting and read classification is paramount — which was also a motive for evolving Kraken 1 into KrakenUniq [6] after the same research group had already authored Kraken 2 [2].

The deliberate focus on a particular set of genera or species allows for more genomes to be included for the organisms of interest while the resulting Genestrip database may still remain small. For example, the tick-related database `tb` from Section 3.4.1 contains almost twice as many k -mers from the chosen genera than the latest publicly available large database `MicrobialDB` for KrakenUniq. As `tb` is deeper than `MicrobialDB` in this respect, it also tends to achieve better recall when analyzing tick related fastq files without loss of precision. We have presented and discussed the corresponding results as part of Section 3.4.3.

Regarding performance, we have examined *database generation as well as fastq file analysis*: Section 3.5.1 has shown that Genestrip empowers users to flexibly create new databases on regular PCs in minutes to just a few hours. As summarized in Tables 8 and 9, we created databases with different applications in mind where the generation of `tb` took the longest with still less than 9 hours while main memory consumption remained below 32 GB in all the cases. Let us contrast Genestrip’s database generation performance with the corresponding resources required under KrakenUniq: As mentioned in Section 1.1, [8] stated that generating a full bacterial database for Kraken 1 and so equivalently for KrakenUniq from the RefSeq “required ca. 2.5 TB of RAM and ca. 11 days” on a suitable computer. However, since then, the RefSeq’s number of prokaryotic genomes alone has increased up by a factor of 5.8 to over 48 Mio. [9]. The additional growth of the RefSeq is also a likely reason why `MicrobialDB` has not been updated anymore since 2023 because ultralarge main

Talk about results on simulated data here.

Relate to results on simulated tick data here.

Wait for final results for the following to rewrite:

memory and extremely high computational power would be required for this. Overall, this development suggests that k -mer databases with very broad and comprehensive coverage (such as by including *all* genomes from the RefSeq category “Bacteria”) have reached a “dead-end” — at least when a lossless storing of k -mers is desired.

In Section 3.5.2 we examined Genestrip’s analysis performance: When the database structure under KrakenUniq and Genestrip was similar then the performance was similar too with KrakenUniq offering about 33% better runtimes. However, when the use case for Genestrip’s small databases was met, then the database size played off with regard to the analysis time: E.g., when analyzing tick-borne bacterial pathogens via the database `tb`, Genestrip was more than ten times faster than KrakenUniq because on a regular PC, KrakenUniq’s large database `MicrobialDB` has to be processed in chunks over several iterations. In a similar scenario where related fastq files were jointly analyzed, Genestrip was over seventy times faster than KrakenUniq because Genestrip spared reloading the database for each fastq file. Regarding main memory, Genestrip’s analysis performance is dominated by the size of the database itself. E.g., it remained below 9 GB of RAM under the database `tb` which has a disk size of about 6 GB.

In use cases where metagenomics is meant to be performed with very few restrictions on what organisms to look out for, then Genestrip may not be the right choice. Although it is certainly possible to create large databases under Genestrip, one would end up with similar challenges as with KrakenUniq (see above). If exactness of results is required in addition to broad coverage, then a two-stage analysis procedure might prove useful: In stage one, a tool such as Kraken 2 with a compressed database for broad coverage may be applied to identify candidate organisms of interest. In stage two, Genestrip might be used to confirm and “drill down” the initial findings by potentially first generating and then using a dedicated database.

5 Availability and requirements

Project name: Genestrip

Project home page: <https://github.com/pfeiferd/genestrip>

Operating system(s): Platform independent

Programming language: Java

Execution requirements: JRE 11 or higher

Build requirements: JDK 11 or higher and Maven 2 or 3

License: Apache 2.0 subject to the Commons Clause License Condition v1.0

Any restrictions to use by non-academics: License needed for commercial use

6 List of abbreviations

t	Placeholder for taxon
f	Placeholder for fastq file
d	Placeholder for k -mer database
$k_{t,f,d}$	Number of k -mers found for taxon t in fastq file f via database d
$u_{t,f,d}$	Number of unique k -mers found for taxon t in fastq file f via database d
$u_{t,d}$	Number of k -mers stored under taxon t in database d
cv_ku	Complete viral database for KrakenUniq
cv_g	Complete viral database for Genestrip
hv_ku	Human virus database for KrakenUniq
hv_g	Human virus database for Genestrip
FPP	False positive probability of a probabilistic filter
BPs	Base pairs
MicrobialDB	Large, comprehensive microbial database for KrakenUniq
tb	Tick-borne bacterial pathogens database for Genestrip
DB	Database

7 Declarations

7.1 Ethics approval and consent to participate

Not applicable

7.2 Consent for publication

Not applicable

7.3 Availability of data and materials

The exact code base for this publication is available under

<https://github.com/pfeiferd/genestrip/releases/tag/v2.5>.

The code base used to generate all of the results from this publication is available under:

<https://github.com/pfeiferd/genestrip-db-exp/tag/v1.7>.

Following [41], the human saliva fastq files were acquired from the Sequence Read Archive (SRA) under the entries: ERX1462737, ERX1462740, SRX2830683, SRX2830684, and SRX2830689. Following [13], the tick-related fastq files were acquired from the Sequence Read Archive (SRA) under the entries SRX13458610, SRX13458622, SRX13458624, SRX13458626, SRX13458627, SRX13458628, SRX13458611 and SRX13458612.

7.4 Competing interests

The authors declare that they have no competing interests.

7.5 Funding

The work behind this publication was not funded.

7.6 Authors' contributions

DP developed the software, designed the evaluation and wrote the manuscript. CR conducted the initial performance experiments and compiled the respective results. MG conducted extended performance experiments and helped to revise the manuscript.

7.7 Acknowledgements

We thank Dr. Richard Zowalla, Simon Schweizer and Jochen Schäfer for providing server-based environments enabling the development and testing of Genestrip. We thank Dr. Robert Friedrich, Dr. Alexander Windberger for their feedback when initially applying the software. Many thanks also to Dr. Alois Heinz for proofreading parts of the manuscript.

References

- [1] Wood DE, Salzberg SL. Kraken: ultrafast metagenomic sequence classification using exact alignments. *Genome Biology*. 2014;15(3):R46. <https://doi.org/10.1186/gb-2014-15-3-r46>.
- [2] Wood DE, Lu J, Langmead B. Improved metagenomic analysis with Kraken 2. *Genome Biology*. 2019;20(1):257. <https://doi.org/10.1186/s13059-019-1891-0>.
- [3] Breitwieser FP, Baker DN, Salzberg SL. KrakenUniq: confident and fast metagenomics classification using unique k-mer counts. *Genome Biology*. 2018;19(1):198. <https://doi.org/10.1186/s13059-018-1568-0>.
- [4] Piro VC, Dadi TH, Seiler E, Reinert K, Renard BY. ganon: precise metagenomics classification against large and up-to-date sets of reference sequences. *Bioinformatics*. 2020 Jul;36(Suppl_1):i12–i20.
- [5] Piro V, Reinert K. ganon2: up-to-date and scalable metagenomics analysis. *NAR Genomics and Bioinformatics*. 2025 07;7(3):lqaf094. <https://doi.org/10.1093/nargab/lqaf094>. <https://academic.oup.com/nargab/article-pdf/7/3/lqaf094/63781699/lqaf094.pdf>.
- [6] Pockrandt C, Zimin AV, Salzberg SL. Metagenomic classification with KrakenUniq on low-memory computers. *J Open Source Softw*. 2022;7(80). <https://doi.org/10.21105/joss.04908>.
- [7] O’Leary NA, Wright MW, Brister JR, Ciufu S, Haddad D, McVeigh R, et al. Reference sequence (RefSeq) database at NCBI: current status, taxonomic expansion, and functional annotation. *Nucleic Acids Res*. 2016 Jan;44(D1):D733–45. <https://doi.org/10.1093/nar/gkv1189>.
- [8] Nasko DJ, Koren S, Phillippy AM, Treangen TJ. RefSeq database growth influences the accuracy of k-mer-based lowest common ancestor species identification. *Genome Biology*. 2018;19(1):165. <https://doi.org/10.1186/s13059-018-1554-6>.
- [9] Goldfarb T, Kodali V, Pujar S, Brover V, Robbertse B, Farrell C, et al. NCBI RefSeq: reference sequence standards through 25 years of curation and annotation. *Nucleic Acids Research*. 2024 11;53(D1):D243–D257. <https://doi.org/10.1093/nar/gkae1038>. <https://academic.oup.com/nar/article-pdf/53/D1/D243/60586509/gkae1038.pdf>.
- [10] Sayers EW, Cavanaugh M, Clark K, Ostell J, Pruitt KD, Karsch-Mizrachi I. GenBank. *Nucleic Acids Research*. 2019 10;48(D1):D84–D86. <https://doi.org/10.1093/nar/gkz956>. <https://academic.oup.com/nar/article-pdf/48/D1/D84/31697602/gkz956.pdf>.

- [11] de la Cuesta-Zuluaga J, Ley RE, Youngblut ND. Struo: a pipeline for building custom databases for common metagenome profilers. *Bioinformatics*. 2020 Apr;36(7):2314–2315. <https://doi.org/10.1093/bioinformatics/btz899>.
- [12] Youngblut ND, Ley RE. Struo2: efficient metagenome profiling database construction for ever-expanding microbial genome datasets. *PeerJ*. 2021;9:e12198. <https://doi.org/10.7717/peerj.12198>.
- [13] Kipp EJ, Lindsey LL, Khoo B, Faulk C, Oliver JD, Larsen PA. Metagenomic surveillance for bacterial tick-borne pathogens using nanopore adaptive sampling. *Scientific Reports*. 2023;13(1):10991. <https://doi.org/10.1038/s41598-023-37134-9>.
- [14] Feldman SI. Make – a program for maintaining computer programs. *Software: Practice and experience*. 1979;9(4):255–265.
- [15] Schoch CL, Ciufo S, Domrachev M, Hottot CL, Kannan S, Khovanskaya R, et al. NCBI Taxonomy: a comprehensive update on curation, resources and tools. *Database (Oxford)*. 2020 Jan;2020. <https://doi.org/10.1093/database/baaa062>.
- [16] Bloom BH. Space/time trade-offs in hash coding with allowable errors. *Commun ACM*. 1970 Jul;13(7):422–426. <https://doi.org/10.1145/362686.362692>.
- [17] Graf TM, Lemire D. Xor Filters: Faster and Smaller Than Bloom and Cuckoo Filters. *ACM J Exp Algorithmics*. 2020 Mar;25. <https://doi.org/10.1145/3376122>.
- [18] Linuwih BP, Satria GB, Mugitama SA, Maulana MS. Space-Efficient Probabilistic Data Structure Ribbon Filter: Analysis, Design, and Optimized Implementation. In: 2022 International Seminar on Application for Technology of Information and Communication (iSemantic); 2022. p. 319–324.
- [19] Breitwieser FP, Pertea M, Zimin AV, Salzberg SL. Human contamination in bacterial genomes has created thousands of spurious proteins. *Genome Res*. 2019 Jun;29(6):954–960. <https://doi.org/10.1101/gr.245373.118>.
- [20] Lu J, Salzberg SL. Removing contaminants from databases of draft genomes. *PLoS Comput Biol*. 2018 Jun;14(6):e1006277. <https://doi.org/10.1371/journal.pcbi.1006277>.
- [21] Morgulis A, Gertz EM, Schäffer AA, Agarwala R. A Fast and Symmetric DUST Implementation to Mask Low-Complexity DNA Sequences. *Journal of Computational Biology*. 2006;13(5):1028–1040. PMID: 16796549. <https://doi.org/10.1089/cmb.2006.13.1028>.
- [22] Delahaye C, Nicolas J. Sequencing DNA with nanopores: Troubles and biases. *PLoS One*. 2021;16(10):e0257521. <https://doi.org/10.1371/journal.pone.0257521>.

- [23] Liu-Wei W, van der Toorn W, Bohn P, Hölzer M, Smyth RP, von Kleist M. Sequencing accuracy and systematic errors of nanopore direct RNA sequencing. *BMC Genomics*. 2024;25(1):528. <https://doi.org/10.1186/s12864-024-10440-w>.
- [24] Wu LY, Wijesekara Y, Piedade GJ, Pappas N, Brussaard CPD, Dutilh BE. Benchmarking bioinformatic virus identification tools using real-world metagenomic data across biomes. *Genome Biology*. 2024;25(1):97. <https://doi.org/10.1186/s13059-024-03236-4>.
- [25] Gourelé H, Karlsson-Lindsjö O, Hayer J, Bongcam-Rudloff E. Simulating Illumina metagenomic data with InSilicoSeq. *Bioinformatics*. 2018 07;35(3):521–522. <https://doi.org/10.1093/bioinformatics/bty630>.
- [26] Katz K, Shutov O, Lapoint R, Kimelman M, Brister J, O'Sullivan C. The Sequence Read Archive: a decade more of explosive growth. *Nucleic Acids Research*. 2021 11;50(D1):D387–D390. <https://doi.org/10.1093/nar/gkab1053>. <https://academic.oup.com/nar/article-pdf/50/D1/D387/42058225/gkab1053.pdf>.
- [27] Hulo C, de Castro E, Masson P, Bougueleret L, Bairoch A, Xenarios I, et al. ViralZone: a knowledge resource to understand virus diversity. *Nucleic Acids Res*. 2011 Jan;39(Database issue):D576–82. <https://doi.org/10.1093/nar/gkq901>.
- [28] Huggins LG, Namgyel U, Wangchuk P, Atapattu U, Traub R, Colella V. Metabarcoding using nanopore sequencing enables identification of diverse and zoonotic vector-borne pathogens from neglected regions: A case study investigating dogs from Bhutan. *One Health*. 2024;19:100839. <https://doi.org/https://doi.org/10.1016/j.onehlt.2024.100839>.
- [29] Miao Y, Guo W, Zhang W, Chen Z, Mian D, Li R, et al. Detection of *Bartonella* spp. in farmed deer (*Artiodactyla*: *Cervidae*) using multiplex assays in the Qinghai-Tibet Plateau, China. *Microbiol Spectr*. 2024 Jul;12(7):e0412023. <https://doi.org/10.1128/spectrum.04120-23>.
- [30] Huggins LG, Colella V, Young ND, Traub RJ. Metabarcoding using nanopore long-read sequencing for the unbiased characterization of apicomplexan haemoparasites. *Mol Ecol Resour*. 2024 Feb;24(2):e13878. <https://doi.org/10.1111/1755-0998.13878>.
- [31] Langmead B.: Kraken 2, KrakenUniq and Bracken indexes. <https://benlangmead.github.io/aws-indexes/k2>.
- [32] Eisen RJ, Eisen L. The Blacklegged Tick, *Ixodes scapularis*: An Increasing Public Health Concern [doi: 10.1016/j.pt.2017.12.006]. *Trends in Parasitology*. 2018 2025/06/22;34(4):295–309. <https://doi.org/10.1016/j.pt.2017.12.006>.

- [33] Johnson TL, Graham CB, Maes SE, Hojgaard A, Fleshman A, Boegler KA, et al. Prevalence and distribution of seven human pathogens in host-seeking *Ixodes scapularis* (Acari: Ixodidae) nymphs in Minnesota, USA. *Ticks and Tick-borne Diseases*. 2018;9(6):1499–1507. <https://doi.org/https://doi.org/10.1016/j.ttbdis.2018.07.009>.
- [34] Janecek E, Mietze A, Goethe R, Schnieder T, Strube C. *Bartonella* spp. infection rate and *B. grahamii* in ticks. *Emerg Infect Dis*. 2012 Oct;18(10):1689–1690. <https://doi.org/10.3201/eid1810.120390>.
- [35] Thomas RE, Karstens RH, Schwan TG. Experimental infection of *Ornithodoros* spp. ticks (Acari: Argasidae) with *Yersinia pestis*. *J Med Entomol*. 1990 Jul;27(4):720–723. <https://doi.org/10.1093/jmedent/27.4.720>.
- [36] Shi H, Li G, Li D, Zhai H, Ji S, Hu Y, et al. Molecular investigation reveals three hemotropic mycoplasmas in cats and three tick species in China. *Frontiers in Veterinary Science*. 2025;Volume 12 - 2025. <https://doi.org/10.3389/fvets.2025.1522904>.
- [37] Yang C, Chu J, Warren RL, Birol I. NanoSim: nanopore sequence read simulator based on statistical characterization. *GigaScience*. 2017 02;6(4):gix010. <https://doi.org/10.1093/gigascience/gix010>.
- [38] Yang C, Lo T, Nip KM, Hafezqorani S, Warren RL, Birol I. Characterization and simulation of metagenomic nanopore sequencing data with Meta-NanoSim. *GigaScience*. 2023 03;12:giad013. <https://doi.org/10.1093/gigascience/giad013>.
- [39] Mehlhorn H. *Die Parasiten des Menschen: Erkrankungen erkennen, bekämpfen und vorbeugen*. SpringerLink: Bücher. Heidelberg: Springer Spektrum; 2023.
- [40] Hillebrand W, Lorenz D, Louis F. *Rebschutz-Taschenbuch*. Heidelberg: Fachverlag Fraund; 1998.
- [41] Samson CA, Whitford W, Snell RG, Jacobsen JC, Lehnert K. Contaminating DNA in human saliva alters the detection of variants from whole genome sequencing. *Sci Rep*. 2020 Nov;10(1):19255. <https://doi.org/10.1038/s41598-020-76022-4>.
- [42] Wensel CR, Pluznick JL, Salzberg SL, Sears CL. Next-generation sequencing: insights to advance clinical investigations of the microbiome. *J Clin Invest*. 2022 Apr;132(7). <https://doi.org/10.1172/JCI154944>.
- [43] Johnson JS, Spakowicz DJ, Hong BY, Petersen LM, Demkowicz P, Chen L, et al. Evaluation of 16S rRNA gene sequencing for species and strain-level microbiome analysis. *Nature Communications*. 2019;10(1):5029. <https://doi.org/10.1038/s41467-019-13036-1>.

- [44] Mendelson AF, Zuluaga MA, Hutton BF, Ourselin S.: What is the distribution of the number of unique original items in a bootstrap sample? Available from: <https://arxiv.org/abs/1602.05822>.
- [45] Sayers EW, Beck J, Bolton EE, Brister JR, Chan J, Connor R, et al. Database resources of the National Center for Biotechnology Information in 2025. *Nucleic Acids Res.* 2025 Jan;53(D1):D20–D29. <https://doi.org/10.1093/nar/gkae979>.
- [46] Chu J, Sadeghi S, Raymond A, Jackman SD, Nip KM, Mar R, et al. BioBloom tools: fast, accurate and memory-efficient host species sequence screening using bloom filters. *Bioinformatics.* 2014 08;30(23):3402–3404. <https://doi.org/10.1093/bioinformatics/btu558>. https://academic.oup.com/bioinformatics/article-pdf/30/23/3402/48931434/bioinformatics_30_23_3402.pdf.
- [47] Flajolet P, Fusy E, Gandouet O, Meunier F. Hyperloglog: The analysis of a near-optimal cardinality estimation algorithm. In: IN AOFA '07: PROCEEDINGS OF THE 2007 INTERNATIONAL CONFERENCE ON ANALYSIS OF ALGORITHMS; 2007. .
- [48] Tange O. GNU Parallel - The Command-Line Power Tool. ;login: The USENIX Magazine. 2011 Feb;36(1):42–47. <https://doi.org/10.5281/zenodo.16303>.
- [49] OEIS Foundation Inc .: The On-Line Encyclopedia of Integer Sequences, Sequence A000045. Accessed: 2026-02-14. <http://oeis.org/A000045>.